

Data606FinalProj

Mathew Katz

2022-12-06

Abstract:

This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective of the dataset is to diagnostically predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset. In 2019, 37.3 million Americans, or 11.3% of the population, had diabetes. 1.4 million Americans are diagnosed with diabetes every year. It is important to be able to diagnose such a worldwide problem. Several constraints were placed on the selection of these instances from a larger database. In particular, all patients here are females at least 21 years old of Pima Indian heritage.

Import Libraries we will need:

```
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'  
  
## The following objects are masked from 'package:stats':  
##  
##   filter, lag  
  
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

```
library(readr)  
library(ggplot2)  
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
library(randomForest)
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##  
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':  
##  
##     margin
```

```
## The following object is masked from 'package:dplyr':  
##  
##     combine
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
library(party)
```

```
## Loading required package: grid
```

```
## Loading required package: mvtnorm
```

```
## Loading required package: modeltools
```

```
## Loading required package: stats4
```

```
## Loading required package: strucchange
```

```
## Loading required package: zoo
```

```
##  
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':  
##  
##     as.Date, as.Date.numeric
```

```
## Loading required package: sandwich
```

```
library(e1071)
```

Read in our Diabetes Dataset:

```
df <- read_csv("diabetes.csv", show_col_types = FALSE)
```

Look at every column in the data frame:

```
df %>% str()
```

```
## spc_tbl_ [768 x 9] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ Pregnancies      : num [1:768] 6 1 8 1 0 5 3 10 2 8 ...
## $ Glucose          : num [1:768] 148 85 183 89 137 116 78 115 197 125 ...
## $ BloodPressure    : num [1:768] 72 66 64 66 40 74 50 0 70 96 ...
## $ SkinThickness    : num [1:768] 35 29 0 23 35 0 32 0 45 0 ...
## $ Insulin          : num [1:768] 0 0 0 94 168 0 88 0 543 0 ...
## $ BMI              : num [1:768] 33.6 26.6 23.3 28.1 43.1 25.6 31 35.3 30.5 0 ...
## $ DiabetesPedigreeFunction: num [1:768] 0.627 0.351 0.672 0.167 2.288 ...
## $ Age              : num [1:768] 50 31 32 21 33 30 26 29 53 54 ...
## $ Outcome          : num [1:768] 1 0 1 0 1 0 1 0 1 1 ...
## - attr(*, "spec")=
## .. cols(
## ..   Pregnancies = col_double(),
## ..   Glucose = col_double(),
## ..   BloodPressure = col_double(),
## ..   SkinThickness = col_double(),
## ..   Insulin = col_double(),
## ..   BMI = col_double(),
## ..   DiabetesPedigreeFunction = col_double(),
## ..   Age = col_double(),
## ..   Outcome = col_double()
## .. )
## - attr(*, "problems")=<externalptr>
```

Based off of pregnancies, glucose level, blood pressure, skin thickness, insulin level, BMI ,Diabetes Pedigree Function, and age let's predict whether someone has diabetes or not but first lets take a quick look at the data:

```
df%>% head()
```

```
## # A tibble: 6 x 9
##   Pregnancies Glucose BloodPressure SkinTh~1 Insulin   BMI Diabe~2   Age Outcome
##   <dbl>      <dbl>         <dbl>    <dbl>    <dbl> <dbl>  <dbl>  <dbl>
## 1         6      148           72      35      0  33.6  0.627   50      1
## 2         1       85           66      29      0  26.6  0.351   31      0
## 3         8      183           64       0      0  23.3  0.672   32      1
## 4         1       89           66      23     94  28.1  0.167   21      0
## 5         0      137           40      35    168  43.1  2.29    33      1
## 6         5      116           74       0      0  25.6  0.201   30      0
## # ... with abbreviated variable names 1: SkinThickness,
## #   2: DiabetesPedigreeFunction
```

The outcome column needs to be changed from a number to a factor (also called categorical or enumerative):

```
df$Outcome <- as.factor(df$Outcome)
levels(df$Outcome) <- c("No", "Yes")
```

Summarize variables:

```
df %>% summary()
```


```
## Pregnancies      Glucose      BloodPressure      SkinThickness
## Min.   : 0.000    Min.   : 0.0    Min.   : 0.00    Min.   : 0.00
## 1st Qu.: 1.000    1st Qu.: 99.0    1st Qu.: 62.00    1st Qu.: 0.00
## Median : 3.000    Median :117.0    Median : 72.00    Median :23.00
## Mean   : 3.845    Mean   :120.9    Mean   : 69.11    Mean   :20.54
## 3rd Qu.: 6.000    3rd Qu.:140.2    3rd Qu.: 80.00    3rd Qu.:32.00
## Max.   :17.000    Max.   :199.0    Max.   :122.00    Max.   :99.00
## Insulin      BMI      DiabetesPedigreeFunction      Age
## Min.   : 0.0    Min.   : 0.00    Min.   :0.0780    Min.   :21.00
## 1st Qu.: 0.0    1st Qu.:27.30    1st Qu.:0.2437    1st Qu.:24.00
## Median : 30.5    Median :32.00    Median :0.3725    Median :29.00
## Mean   : 79.8    Mean   :31.99    Mean   :0.4719    Mean   :33.24
## 3rd Qu.:127.2    3rd Qu.:36.60    3rd Qu.:0.6262    3rd Qu.:41.00
## Max.   :846.0    Max.   :67.10    Max.   :2.4200    Max.   :81.00
## Outcome
## No :500
## Yes:268
##
##
##
##
```

Plots:

```
plotpreg <- ggplot(data = df, aes(x = Pregnancies)) +
  geom_histogram(color = "black", fill = "red") +
  labs(title = "Pregnancies Histogram Plot", x = "Pregnancies", y = "Count")
plotpreg
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```


Pregnancies Histogram Plot



```
plotgluc <- ggplot(data = df, aes(x = Glucose)) +  
  geom_histogram(color = "black", fill = "orange") +  
  labs(title = "Glucose Histogram Plot", x = "Glucose", y = "Count")  
plotgluc
```


```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

Glucose Histogram Plot



```
plotbp <- ggplot(data = df, aes(x = BloodPressure)) +  
  geom_histogram(color = "black", fill = "yellow") +  
  labs(title = "Blood Pressure Histogram Plot", x = "Blood Pressure", y = "Count")  
plotbp
```

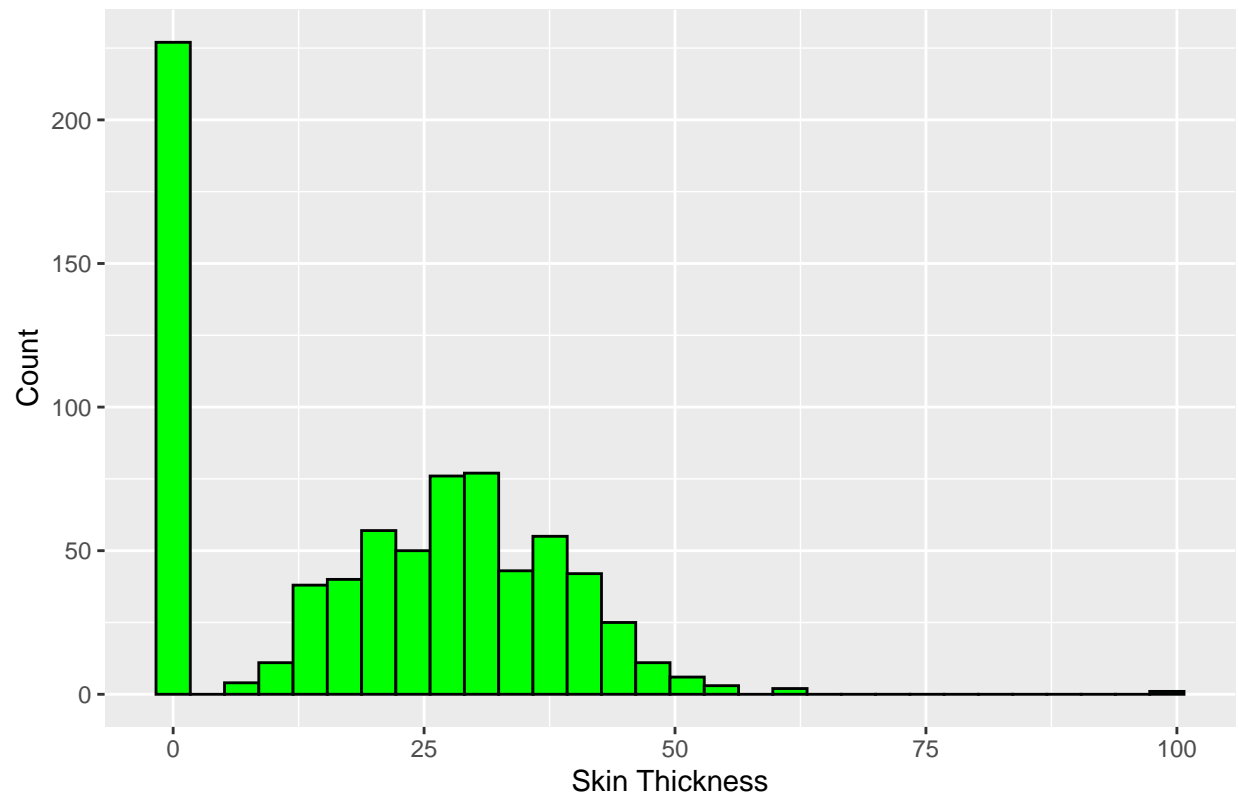
```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



```
plotst <- ggplot(data = df, aes(x = SkinThickness)) +  
  geom_histogram(color = "black", fill = "green") +  
  labs(title = "Skin Thickness Histogram Plot", x = "Skin Thickness", y = "Count")  
plotst
```

'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.


Skin Thickness Histogram Plot



```
plotins <- ggplot(data = df, aes(x = Insulin)) +  
  geom_histogram(color = "black", fill = "blue") +  
  labs(title = "Insulin Histogram Plot", x = "Insulin", y = "Count")  
plotins
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



Insulin Histogram Plot



```
plotbmi <- ggplot(data = df, aes(x = BMI)) +  
  geom_histogram(color = "black", fill = "purple") +  
  labs(title = "BMI Plot", x = "BMI", y = "Count")  
plotbmi
```

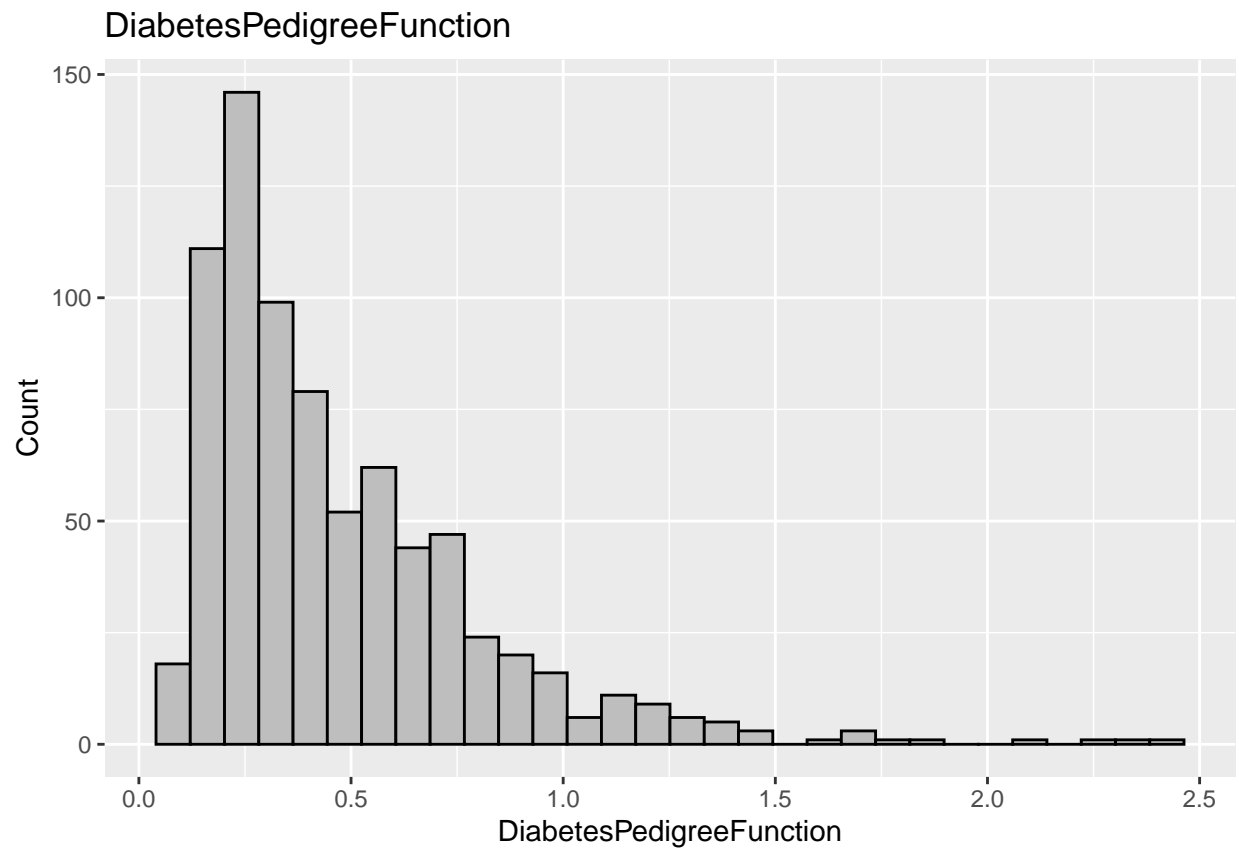
```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

BMI Plot




```
plotdpf <- ggplot(data = df, aes(x = DiabetesPedigreeFunction)) +  
  geom_histogram(color = "black", fill = "grey") +  
  labs(title = "DiabetesPedigreeFunction", x = "DiabetesPedigreeFunction", y = "Count")  
plotdpf
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

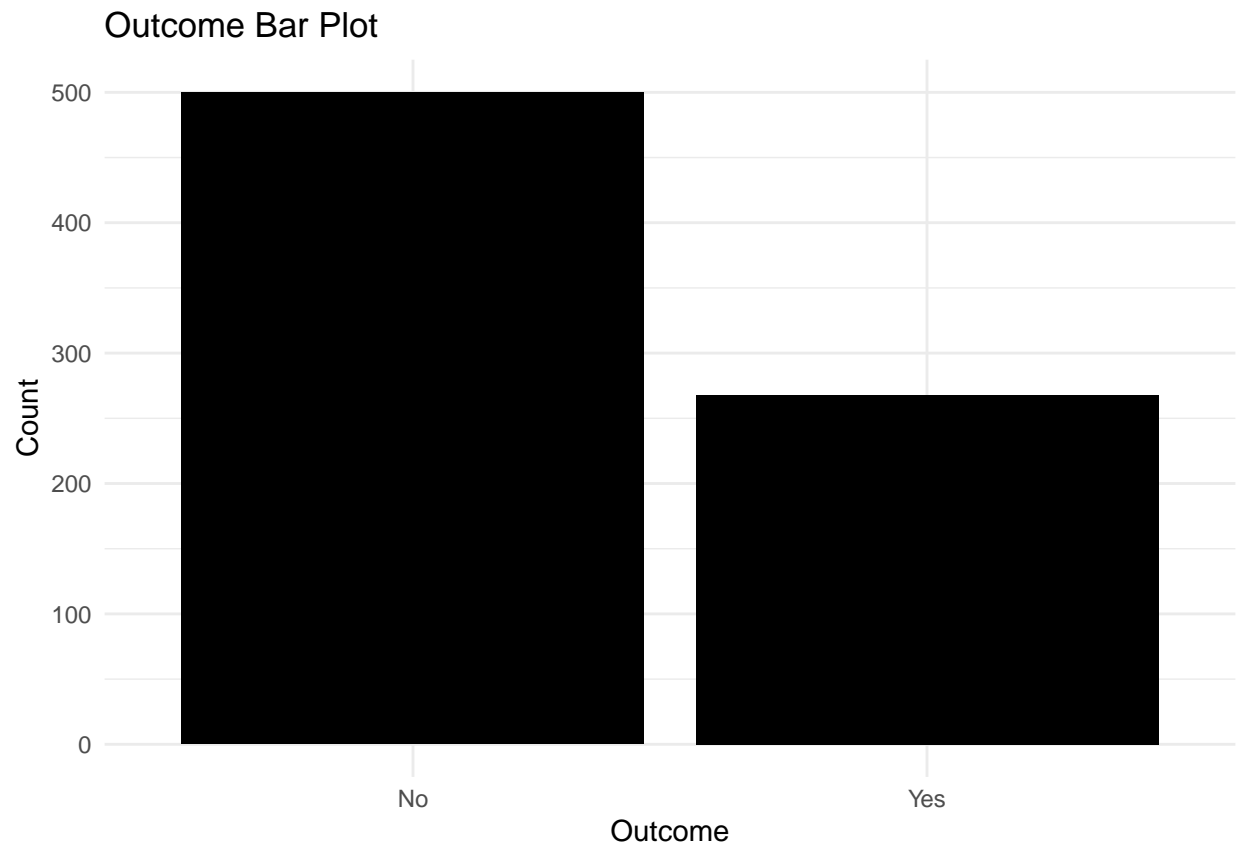


```
plotage <- ggplot(data = df, aes(x = Age)) +  
  geom_histogram(color = "black", fill = "white") +  
  labs(title = "Age", x = "Age", y = "Count")  
plotage
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```




```
plotout <- ggplot(data = df, aes(x = Outcome)) +  
  stat_count(fill = "black") +  
  labs(title = "Outcome Bar Plot", x = "Outcome", y = "Count") +  
  theme_minimal()  
plotout
```



Correlation of Variables:

```
df_corr <- df[-9]
df_corr <- cor(df_corr)
corrplot(df_corr, method = "color", type = "lower",
          addCoef.col = "black", col = COL2("RdYlBu"), number.cex = .5, tl.cex = .5)
```



There are moderately positive correlations between the Age and Pregnancy, and the Insulin and Skin Thickness attributes. This indicates that as the age of the patients increased so did the number of pregnancies, also as the quantity of insulin administered to the patients increased; the skin thickness increased likewise.

Weak positive correlations can also be observed in the following attributes of the dataset; Insulin & Glucose, BMI & Skin Thickness, Blood Pressure & BMI, Age & Blood Pressure...

Time to Train-Test-Split:

```
nrows <- NROW(df)
set.seed(42)
index <- sample(1:nrows, 0.7 * nrows)
train <- df[index,]
test <- df[-index,]
```

Random Forest Classifier:

```
learn_rf <- randomForest(Outcome~., data=train, ntree=500, proximity=T, importance=T)
pre_rf <- predict(learn_rf, test[, -9])
cm_rf <- confusionMatrix(pre_rf, test$Outcome)
```

```
cm_rf
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 134  28
##           Yes  22  47
##
##           Accuracy : 0.7835
##           95% CI : (0.7248, 0.8349)
##           No Information Rate : 0.6753
##           P-Value [Acc > NIR] : 0.0001885
##
##           Kappa : 0.4959
##
##  Mcnemar's Test P-Value : 0.4795001
##
##           Sensitivity : 0.8590
##           Specificity : 0.6267
##           Pos Pred Value : 0.8272
##           Neg Pred Value : 0.6812
##           Prevalence : 0.6753
##           Detection Rate : 0.5801
##           Detection Prevalence : 0.7013
##           Balanced Accuracy : 0.7428
##
##           'Positive' Class : No
##
```

CTree Classifier:

```
learn_ct <- ctree(Outcome~., data=train, controls=ctree_control(maxdepth=2))

pre_ct   <- predict(learn_ct, test[, -9])

cm_ct    <- confusionMatrix(pre_ct, test$Outcome)

cm_ct
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  No  Yes
##           No 132  36
##           Yes  24  39
##
##           Accuracy : 0.7403
##           95% CI : (0.6787, 0.7956)
##           No Information Rate : 0.6753
##           P-Value [Acc > NIR] : 0.01935
##
```

```
##                Kappa : 0.382
##
## Mcnemar's Test P-Value : 0.15558
##
##          Sensitivity : 0.8462
##          Specificity : 0.5200
##          Pos Pred Value : 0.7857
##          Neg Pred Value : 0.6190
##          Prevalence : 0.6753
##          Detection Rate : 0.5714
##          Detection Prevalence : 0.7273
##          Balanced Accuracy : 0.6831
##
##          'Positive' Class : No
##
```

Naive Bayes Classifier:

```
learn_nb <- naiveBayes(train[,-9], train$Outcome)

pre_nb <- predict(learn_nb, test[,-9])

cm_nb <- confusionMatrix(pre_nb, test$Outcome)

cm_nb
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction  No Yes
##          No 125 27
##          Yes 31 48
##
##          Accuracy : 0.7489
##          95% CI : (0.6878, 0.8035)
##          No Information Rate : 0.6753
##          P-Value [Acc > NIR] : 0.009121
##
##          Kappa : 0.4353
##
## Mcnemar's Test P-Value : 0.693641
##
##          Sensitivity : 0.8013
##          Specificity : 0.6400
##          Pos Pred Value : 0.8224
##          Neg Pred Value : 0.6076
##          Prevalence : 0.6753
##          Detection Rate : 0.5411
##          Detection Prevalence : 0.6580
##          Balanced Accuracy : 0.7206
##
##          'Positive' Class : No
##
```


The Random Forest Classifier performed the best with a 78% accuracy, 85.9% Sensitivity, and 63% Specificity. Let's see how our model does on two test subjects. A yes and a no diabetes:

```
Y <- test[1,]
N <- test[2,]

print(Y)
```

```
## # A tibble: 1 x 9
##   Pregnancies Glucose BloodPressure SkinTh~1 Insulin   BMI Diabe~2   Age Outcome
##         <dbl>   <dbl>         <dbl>   <dbl>   <dbl> <dbl>   <dbl> <dbl> <fct>
## 1           6    148           72     35     0  33.6   0.627    50 Yes
## # ... with abbreviated variable names 1: SkinThickness,
## #   2: DiabetesPedigreeFunction
```

```
print(N)
```

```
## # A tibble: 1 x 9
##   Pregnancies Glucose BloodPressure SkinTh~1 Insulin   BMI Diabe~2   Age Outcome
##         <dbl>   <dbl>         <dbl>   <dbl>   <dbl> <dbl>   <dbl> <dbl> <fct>
## 1           1     89           66     23    94  28.1   0.167    21 No
## # ... with abbreviated variable names 1: SkinThickness,
## #   2: DiabetesPedigreeFunction
```

Remove the outcome variable so it isn't there to mess with the model prediction:

```
Y$Outcome <- NULL
N$Outcome <- NULL
```

Function that takes in a patient's variables and predicts if they have diabetes or not:

```
patient_diabetes_predict <- function(new, method=learn_rf) {
  new_pre <- predict(method, new)
  new_res <- as.character(new_pre)
  return(paste("Result: ", new_res, sep=""))
}
```

Yes patient's prediction:

```
patient_diabetes_predict(Y)
```

```
## [1] "Result: Yes"
```

No patient's prediction:

```
patient_diabetes_predict(N)
```

```
## [1] "Result: No"
```