

SentAnalysisData607

Mathew Katz

2022-11-06

Firstly I am going to post a bunch of code from Text Mining with R, Chapter 2 (Sentiment Analysis.)
{<https://www.tidytextmining.com/sentiment.html>}

```
library(tidytext)
get_sentiments("afinn")
```

```
## # A tibble: 2,477 x 2
##   word      value
##   <chr>    <dbl>
## 1 abandon     -2
## 2 abandoned   -2
## 3 abandons    -2
## 4 abducted    -2
## 5 abduction   -2
## 6 abductions   -2
## 7 abhor       -3
## 8 abhorred    -3
## 9 abhorrent   -3
## 10 abhors     -3
## # ... with 2,467 more rows
```

```
get_sentiments("bing")
```

```
## # A tibble: 6,786 x 2
##   word      sentiment
##   <chr>    <chr>
## 1 2-faces  negative
## 2 abnormal negative
## 3 abolish negative
## 4 abominable negative
## 5 abominably negative
## 6 abominate negative
## 7 abomination negative
## 8 abort    negative
## 9 aborted  negative
## 10 aborts  negative
## # ... with 6,776 more rows
```

```
get_sentiments("nrc")
```

```
## # A tibble: 13,872 x 2
##   word      sentiment
##   <chr>     <chr>
## 1 abacus    trust
## 2 abandon   fear
## 3 abandon   negative
## 4 abandon   sadness
## 5 abandoned anger
## 6 abandoned fear
## 7 abandoned negative
## 8 abandoned sadness
## 9 abandonment anger
## 10 abandonment fear
## # ... with 13,862 more rows
```

```
library(janeaustrer)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(stringr)

tidy_books <- austen_books() %>%
  group_by(book) %>%
  mutate(
    linenumber = row_number(),
    chapter = cumsum(str_detect(text,
                                regex("^chapter [\\divxlc]",
                                      ignore_case = TRUE)))) %>%
  ungroup() %>%
  unnest_tokens(word, text)
```

```
nrc_joy <- get_sentiments("nrc") %>%
  filter(sentiment == "joy")

tidy_books %>%
  filter(book == "Emma") %>%
  inner_join(nrc_joy) %>%
  count(word, sort = TRUE)
```

```
## Joining, by = "word"
```

```
## # A tibble: 301 x 2
##   word      n
##   <chr>    <int>
## 1 good      359
## 2 friend    166
## 3 hope      143
## 4 happy     125
## 5 love      117
## 6 deal       92
## 7 found      92
## 8 present    89
## 9 kind       82
## 10 happiness 76
## # ... with 291 more rows
```

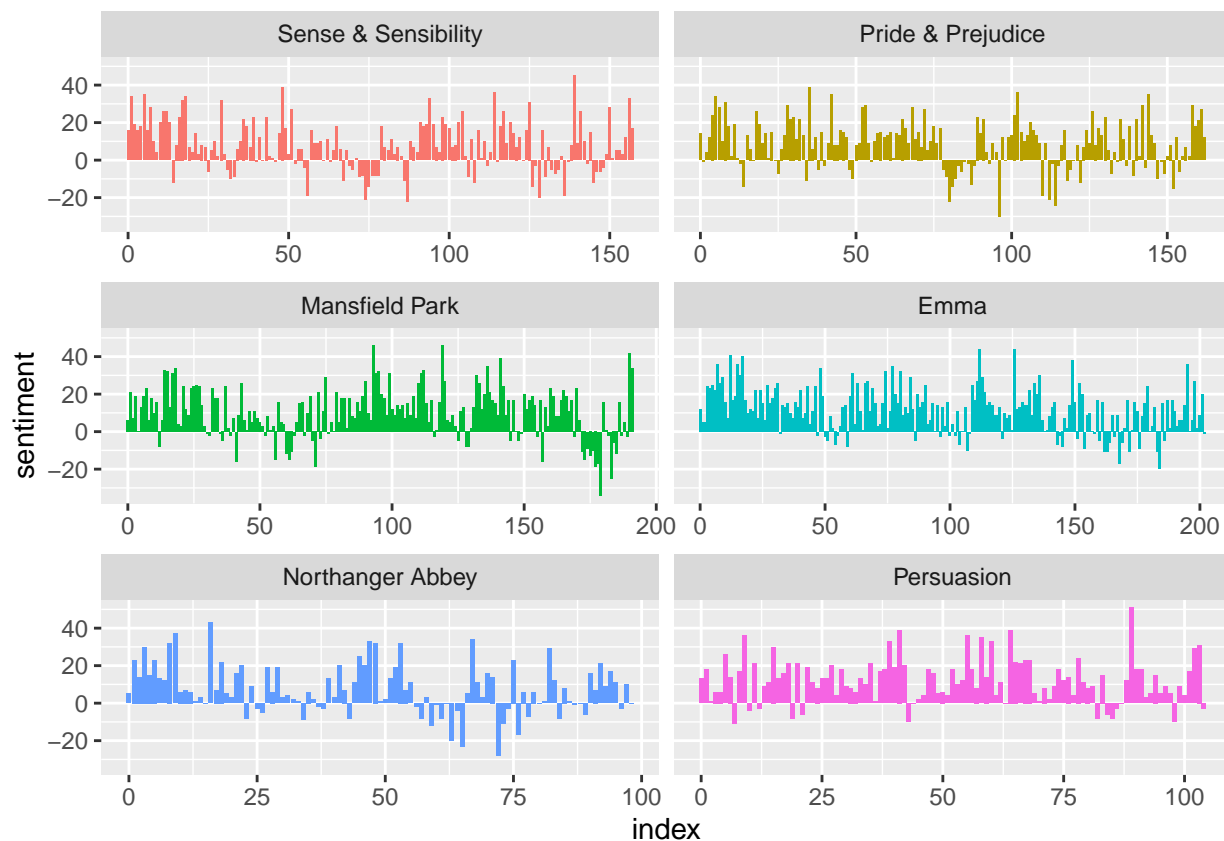
```
library(tidyr)

jane_austen_sentiment <- tidy_books %>%
  inner_join(get_sentiments("bing")) %>%
  count(book, index = linenumber %/% 80, sentiment) %>%
  pivot_wider(names_from = sentiment, values_from = n, values_fill = 0) %>%
  mutate(sentiment = positive - negative)
```

```
## Joining, by = "word"
```

```
library(ggplot2)

ggplot(jane_austen_sentiment, aes(index, sentiment, fill = book)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~book, ncol = 2, scales = "free_x")
```



```
pride_prejudice <- tidy_books %>%
  filter(book == "Pride & Prejudice")
```

```
pride_prejudice
```

```
## # A tibble: 122,204 x 4
##   book          linenumber chapter word
##   <fct>          <int>    <int> <chr>
## 1 Pride & Prejudice      1      0 pride
## 2 Pride & Prejudice      1      0 and
## 3 Pride & Prejudice      1      0 prejudice
## 4 Pride & Prejudice      3      0 by
## 5 Pride & Prejudice      3      0 jane
## 6 Pride & Prejudice      3      0 austen
## 7 Pride & Prejudice      7      1 chapter
## 8 Pride & Prejudice      7      1 1
## 9 Pride & Prejudice     10      1 it
## 10 Pride & Prejudice     10      1 is
## # ... with 122,194 more rows
```

```
afinn <- pride_prejudice %>%
  inner_join(get_sentiments("afinn")) %>%
  group_by(index = linenumber %/% 80) %>%
  summarise(sentiment = sum(value)) %>%
  mutate(method = "AFINN")
```


```
## Joining, by = "word"
```

```
bing_and_nrc <- bind_rows(
  pride_prejudice %>%
    inner_join(get_sentiments("bing")) %>%
    mutate(method = "Bing et al."),
  pride_prejudice %>%
    inner_join(get_sentiments("nrc")) %>%
    filter(sentiment %in% c("positive",
                           "negative"))
) %>%
  mutate(method = "NRC")) %>%
count(method, index = linenumber %/% 80, sentiment) %>%
pivot_wider(names_from = sentiment,
             values_from = n,
             values_fill = 0) %>%
mutate(sentiment = positive - negative)
```

```
## Joining, by = "word"
```

```
## Joining, by = "word"
```

```
bind_rows(afinn,
          bing_and_nrc) %>%
  ggplot(aes(index, sentiment, fill = method)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~method, ncol = 1, scales = "free_y")
```



```
get_sentiments("nrc") %>%
  filter(sentiment %in% c("positive", "negative")) %>%
  count(sentiment)
```

```
## # A tibble: 2 x 2
##   sentiment      n
##   <chr>      <int>
## 1 negative   3316
## 2 positive   2308
```

```
get_sentiments("bing") %>%
  count(sentiment)
```

```
## # A tibble: 2 x 2
##   sentiment      n
##   <chr>      <int>
## 1 negative   4781
## 2 positive   2005
```

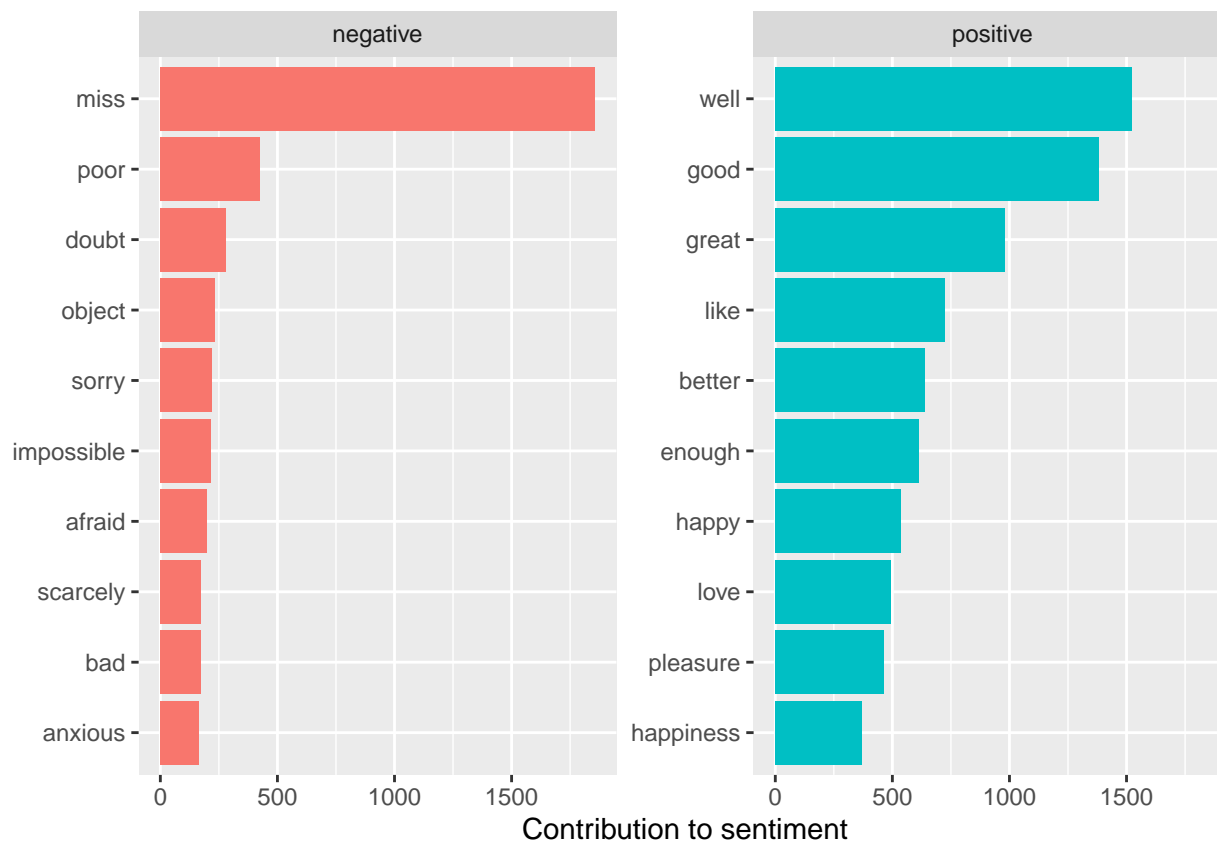
```
bing_word_counts <- tidy_books %>%
  inner_join(get_sentiments("bing")) %>%
  count(word, sentiment, sort = TRUE) %>%
  ungroup()
```

```
## Joining, by = "word"
```

```
bing_word_counts
```

```
## # A tibble: 2,585 x 3
##   word      sentiment      n
##   <chr>      <chr>    <int>
## 1 miss      negative    1855
## 2 well      positive    1523
## 3 good      positive    1380
## 4 great     positive     981
## 5 like      positive     725
## 6 better    positive     639
## 7 enough    positive     613
## 8 happy     positive     534
## 9 love      positive     495
## 10 pleasure positive     462
## # ... with 2,575 more rows
```

```
bing_word_counts %>%
  group_by(sentiment) %>%
  slice_max(n, n = 10) %>%
  ungroup() %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(n, word, fill = sentiment)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~sentiment, scales = "free_y") +
  labs(x = "Contribution to sentiment",
       y = NULL)
```



```
custom_stop_words <- bind_rows(tibble(word = c("miss"),
                                       lexicon = c("custom")),
                               stop_words)
```

```
custom_stop_words
```

```
## # A tibble: 1,150 x 2
##   word      lexicon
##   <chr>    <chr>
## 1 miss    custom
## 2 a       SMART
## 3 a's     SMART
## 4 able    SMART
## 5 about   SMART
## 6 above   SMART
## 7 according SMART
## 8 accordingly SMART
## 9 across  SMART
## 10 actually SMART
## # ... with 1,140 more rows
```

```
library(wordcloud)
```

```
## Loading required package: RColorBrewer
```

```
tidy_books %>%
  anti_join(stop_words) %>%
  count(word) %>%
  with(wordcloud(word, n, max.words = 100))

## Joining, by = "word"

## Warning in wordcloud(word, n, max.words = 100): elizabeth could not be fit on
## page. It will not be plotted.
```



```
library(reshape2)

##
## Attaching package: 'reshape2'

## The following object is masked from 'package:tidyr':
##
##      smiths

tidy_books %>%
  inner_join(get_sentiments("bing")) %>%
  count(word, sentiment, sort = TRUE) %>%
  acast(word ~ sentiment, value.var = "n", fill = 0) %>%
  comparison.cloud(colors = c("gray20", "gray80"),
                   max.words = 100)
```



```
## Joining, by = "word"
```



```
p_and_p_sentences <- tibble(text = prideprejudice) %>%
  unnest_tokens(sentence, text, token = "sentences")
```

```
p_and_p_sentences$sentence[2]
```

```
## [1] "by jane austen"
```

```
austen_chapters <- austen_books() %>%
  group_by(book) %>%
  unnest_tokens(chapter, text, token = "regex",
                 pattern = "Chapter|CHAPTER [\\dIVXLC]") %>%
  ungroup()

austen_chapters %>%
  group_by(book) %>%
  summarise(chapters = n())
```

```
## # A tibble: 6 x 2
##   book                      chapters
##   <fct>                    <int>
## 1 Sense & Sensibility      51
## 2 Pride & Prejudice       62
```

```
## 3 Mansfield Park          49
## 4 Emma                    56
## 5 Northanger Abbey       32
## 6 Persuasion              25
```

```
bingnegative <- get_sentiments("bing") %>%
  filter(sentiment == "negative")

wordcounts <- tidy_books %>%
  group_by(book, chapter) %>%
  summarize(words = n())
```

```
## 'summarise()' has grouped output by 'book'. You can override using the
## '.groups' argument.
```

```
tidy_books %>%
  semi_join(bingnegative) %>%
  group_by(book, chapter) %>%
  summarize(negativewords = n()) %>%
  left_join(wordcounts, by = c("book", "chapter")) %>%
  mutate(ratio = negativewords/words) %>%
  filter(chapter != 0) %>%
  slice_max(ratio, n = 1) %>%
  ungroup()
```

```
## Joining, by = "word"
## 'summarise()' has grouped output by 'book'. You can override using the
## '.groups' argument.
```

```
## # A tibble: 6 x 5
##   book          chapter negativewords words  ratio
##   <fct>         <int>         <int> <int>  <dbl>
## 1 Sense & Sensibility    43          161  3405 0.0473
## 2 Pride & Prejudice     34           111  2104 0.0528
## 3 Mansfield Park       46           173  3685 0.0469
## 4 Emma                 15           151  3340 0.0452
## 5 Northanger Abbey     21           149  2982 0.0500
## 6 Persuasion            4            62  1807 0.0343
```

Let's extend the code in two ways: Work with a different corpus of our choosing, and Incorporate at least one additional sentiment lexicon (possibly from another R package that I've found through research).

```
library('readr')
scripts <- read_csv("RickAndMortyScripts.csv")
```

```
## Rows: 1905 Columns: 6
## -- Column specification -----
## Delimiter: ","
## chr (3): episode name, name, line
## dbl (3): index, season no., episode no.
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
nrc <- get_sentiments("nrc")
bing <- get_sentiments("bing")
afinn <- get_sentiments("afinn")
```

```
scripts = scripts %>% rename(Index = "index",
                             Season.No = "season no.",
                             Episode.No = "episode no.",
                             Episode.Name = "episode name",
                             Character.Name = "name",
                             Dialog = "line")
```

```
# Head of the table
head(scripts, 4)
```

```
## # A tibble: 4 x 6
##   Index Season.No Episode.No Episode.Name Character.Name Dialog
##   <dbl>   <dbl>   <dbl> <chr>         <chr>         <chr>
## 1     0     1     1 Pilot         Rick         Morty! You gotta come ~
## 2     1     1     1 Pilot         Morty        What, Rick? What's goi~
## 3     2     1     1 Pilot         Rick         I got a surprise for y~
## 4     3     1     1 Pilot         Morty        It's the middle of the~
```

```
# Tail of the table
tail(scripts, 4)
```

```
## # A tibble: 4 x 6
##   Index Season.No Episode.No Episode.Name Character.Name Dialog
##   <dbl>   <dbl>   <dbl> <chr>         <chr>         <chr>
## 1  2484     3     7 Tales From the Citadel Rick         Got some of ~
## 2  2485     3     7 Tales From the Citadel Morty        I'm really h~
## 3  2486     3     7 Tales From the Citadel Rick         Pssh! Not at~
## 4  2487     3     7 Tales From the Citadel Morty        Whoo! Yeah! ~
```

```
# Summary
summary(scripts)
```

```
##      Index      Season.No      Episode.No      Episode.Name
## Min.   : 0      Min.   :1.000      Min.   : 1.000      Length:1905
## 1st Qu.: 548    1st Qu.:1.000    1st Qu.: 1.000      Class :character
## Median :1164    Median :2.000    Median : 3.000      Mode  :character
## Mean   :1190    Mean   :2.155    Mean   : 3.208
## 3rd Qu.:1844    3rd Qu.:3.000    3rd Qu.: 5.000
## Max.   :2487    Max.   :3.000    Max.   :10.000
## Character.Name      Dialog
## Length:1905      Length:1905
## Class :character  Class :character
## Mode  :character  Mode  :character
##
##
##
```

Clean Corpus Function: This predefined function is going to clean the text from:

the punctuation - removePunctuation extra white space - stripWhitespace transforms to lower case - tolower
stopwords (common words that should be ignored) - stopwords numbers - removeNumbers

```
cleanCorpus <- function(text){  
  # punctuation, whitespace, lowercase, numbers  
  text.tmp <- tm_map(text, removePunctuation)  
  text.tmp <- tm_map(text.tmp, stripWhitespace)  
  text.tmp <- tm_map(text.tmp, content_transformer(tolower))  
  text.tmp <- tm_map(text.tmp, removeNumbers)  
  
  # removes stopwords  
  stopwords_remove <- c(stopwords("en"), c("thats", "weve", "hes", "theres", "ive", "im",  
                                             "will", "can", "cant", "dont", "youve", "us",  
                                             "youre", "youll", "theyre", "whats", "didnt"))  
  text.tmp <- tm_map(text.tmp, removeWords, stopwords_remove)  
  
  return(text.tmp)  
}
```

These predefined functions will process the text depending on the case:

Unigrams take only 1 word at a time Bigrams take 2 sequential words at a time Trigrams (you guessed) take 3 sequential words at a time Eg. text: "come on morty"

Unigram: "come", "on", "morty" Bigram: "come on", "on morty" Trigram: "come on morty" Term Document Matrix: it's a mathematical matrix that describes the frequency of terms that occur in a collection of documents. More simply put, is a matrix that has on:

rows - words that can be found in the analysed documents columns - the documents in order values - the frequency of each word in each document

Unigram:

```
frequentTerms <- function(text){  
  # create the matrix  
  s.cor <- VCorpus(VectorSource(text))  
  s.cor.cl <- cleanCorpus(s.cor)  
  s.tdm <- TermDocumentMatrix(s.cor.cl)  
  s.tdm <- removeSparseTerms(s.tdm, 0.999)  
  m <- as.matrix(s.tdm)  
  word_freqs <- sort(rowSums(m), decreasing = T)  
  
  # change to dataframe  
  dm <- data.frame(word=names(word_freqs), freq=word_freqs)  
  
  return(dm)  
}
```

Bigram:

```
# Bigram tokenizer  
tokenizer_2 <- function(x){  
  NGramTokenizer(x, Weka_control(min=2, max=2))  
}
```

```

}

# Bigram function
frequentBigrams <- function(text){

  s.cor <- VCorpus(VectorSource(text))
  s.cor.cl <- cleanCorpus(s.cor)
  s.tdm <- TermDocumentMatrix(s.cor.cl, control=list(tokenize=tokenizer_2))
  s.tdm <- removeSparseTerms(s.tdm, 0.999)
  m <- as.matrix(s.tdm)
  word_freqs <- sort(rowSums(m), decreasing=T)
  dm <- data.frame(word=names(word_freqs), freq=word_freqs)

  return(dm)
}

```

Trigram:

```

# Trigram tokenizer
tokenizer_3 <- function(x){
  NGramTokenizer(x, Weka_control(min=3, max=3))
}

# Trigram function
frequentTrigrams <- function(text){

  s.cor <- VCorpus(VectorSource(text))
  s.cor.cl <- cleanCorpus(s.cor)
  s.tdm <- TermDocumentMatrix(s.cor.cl, control=list(tokenize=tokenizer_3))
  s.tdm <- removeSparseTerms(s.tdm, 0.999)
  m <- as.matrix(s.tdm)
  word_freqs <- sort(rowSums(m), decreasing=T)
  dm <- data.frame(word=names(word_freqs), freq=word_freqs)

  return(dm)
}

```

Bing Lexicon categorizes the words into positives and negatives.

To be able to do so in our data, first we make a dataframe that splits all the words in 1 dialogue onto rows. Afterwards, we can join our data with the lexicon, leaving us with a beautiful classification of our words.

```

# Creating our tokens
tokens <- scripts %>%
  mutate(dialogue = as.character(scripts$Dialog)) %>%
  unnest_tokens(word, dialogue)

tokens %>% head(5) %>% select(Character.Name, word)

```

```

## # A tibble: 5 x 2
##   Character.Name word
##   <chr>          <chr>
## 1 Rick          morty

```

```
tokens %>%
  # append the bing sentiment and prepare the data
  inner_join(bing, "word") %>%
  count(word, sentiment, sort=T) %>%
  acast(word ~ sentiment, value.var = "n", fill=0) %>%

# wordcloud
comparison.cloud(colors=c("#991D1D", "#327CDE"), max.words = 100)
```

[illegible]

positive

The nrc lexicon categorizes the words in 10 moods:

```
sentiments <- tokens %>%
  inner_join(nrc, "word") %>%
  count(sentiment, sort=T)
```

```
## # A tibble: 10 x 2
```

```
##      sentiment      n
##      <chr>        <int>
## 1 positive      977
## 2 negative      901
## 3 trust         645
## 4 anticipation  591
## 5 fear          567
## 6 joy           494
## 7 anger         415
## 8 sadness       414
## 9 disgust       312
## 10 surprise     266
```


Afinn Lexicon ranks every word from -5 to 5, where:

-5 being the most negative +5 being the most positive

```
tokens %>%
  # Count how many word per value
  inner_join(afinn, "word") %>%
  count(value, sort=T) %>%

  # Plot
  ggplot(aes(x=value, y=n)) +
  geom_bar(stat="identity", aes(fill=n), show.legend = F, width = 0.5) +
  geom_label(aes(label=n)) +
  scale_fill_gradient(low="#85C1E9", high="#3498DB") +
  scale_x_continuous(breaks=seq(-5, 5, 1)) +
  labs(x="Score", y="Frequency", title="Word count distribution over intensity of sentiment: Neg -> Pos") +
  theme_bw()
```

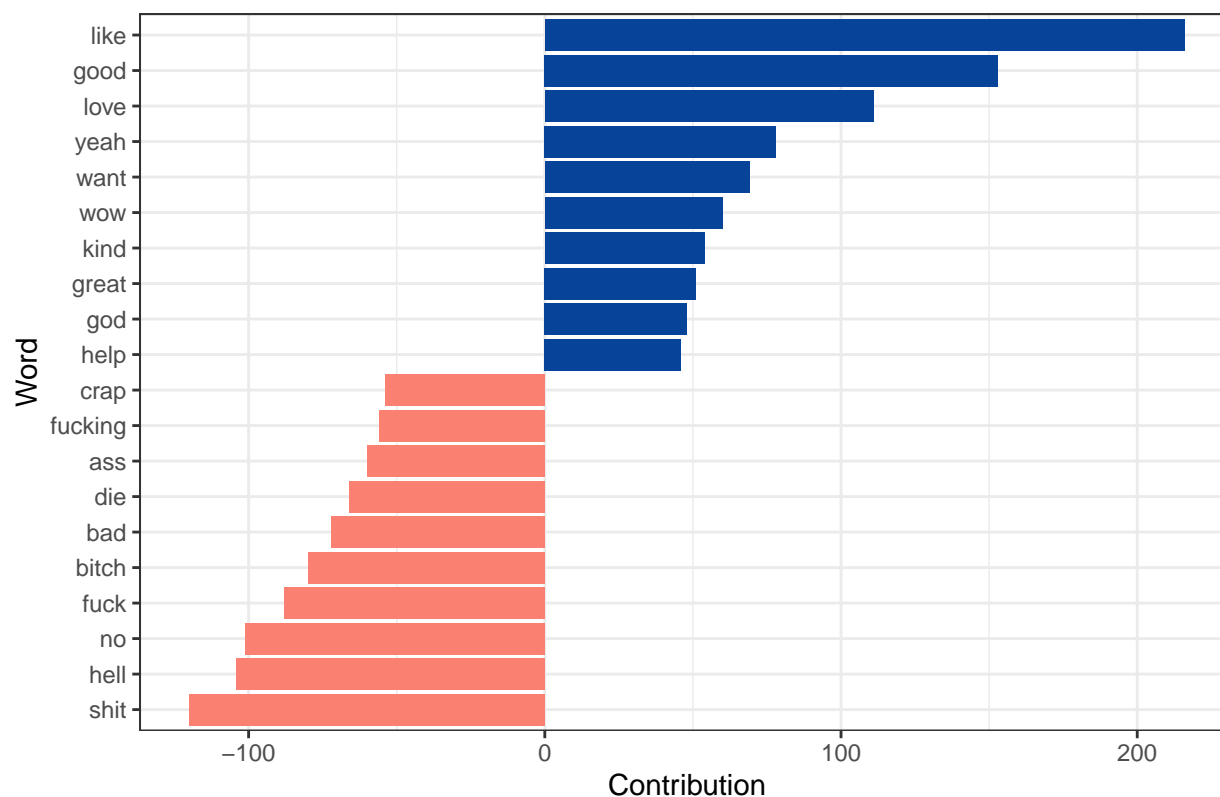
Word count distribution over intensity of sentiment: Neg → Pos



```
tokens %>%
  # by word and value count number of occurrences
  inner_join(afinn, "word") %>%
  count(word, value, sort=T) %>%
  mutate(contribution = n * value,
          sentiment = ifelse(contribution<=0, "Negative", "Positive")) %>% #another variable
  arrange(desc(abs(contribution))) %>%
  head(20) %>%

# plot
ggplot(aes(x=reorder(word, contribution), y=contribution, fill=sentiment)) +
  geom_col(aes(fill=sentiment), show.legend = F) +
  labs(x="Word", y="Contribution", title="Words with biggest contributions in positive/negative moods")
  coord_flip() +
  scale_fill_manual(values=c("#FA8072", "#08439A")) +
  theme_bw()
```


Words with biggest contributions in positive/negative moods



```
library("stopwords")

# Create a dataframe with stopwords
stopwords_script <- tibble(word = c(stopwords("en"), c("thats", "weve", "hes", "theres", "ive", "im",
  "will", "can", "cant", "dont", "youve", "us",
  "youre", "youll", "theyre", "whats", "didnt"))))

print(stopwords_script)
```

```
## # A tibble: 192 x 1
##   word
##   <chr>
## 1 i
## 2 me
## 3 my
## 4 myself
## 5 we
## 6 our
## 7 ours
## 8 ourselves
## 9 you
## 10 your
## # ... with 182 more rows
```

```
# Create the dataframe of tokens
scripts %>%
```

```

mutate(dialogue = as.character(scripts$Dialog)) %>%
filter(Character.Name %in% c("Rick", "Morty", "Beth", "Jerry", "Summer")) %>%

# removes stopwords
unnest_tokens(word, dialogue) %>%
anti_join(stopwords_script, by="word") %>%

# top N frequent words per character
count(Character.Name, word) %>%
group_by(Character.Name) %>%
arrange(desc(n)) %>%
slice(1:10) %>%

mutate(word2 = factor(paste(word, Character.Name, sep="__"),
                      levels = rev(paste(word, Character.Name, sep="__"))))

```

```

## # A tibble: 50 x 4
## # Groups:   Character.Name [5]
##   Character.Name word      n word2
##   <chr>          <chr> <int> <fct>
## 1 Beth          jerry    22 jerry__Beth
## 2 Beth          dad      12 dad__Beth
## 3 Beth          oh       12 oh__Beth
## 4 Beth          summer   12 summer__Beth
## 5 Beth          know     11 know__Beth
## 6 Beth          morty    10 morty__Beth
## 7 Beth          want     10 want__Beth
## 8 Beth          like      9 like__Beth
## 9 Beth          mean      9 mean__Beth
## 10 Beth         get       8 get__Beth
## # ... with 40 more rows

```

Rick and Morty's script made this project very enjoyable.